



From Engineering To Research

The KG Construction Use Case

David Chaves-Fraga, Ontology Engineering Group
Universidad Politécnica de Madrid, Spain

Samaneh Jozashoori, SDM-TIB

Maria-Ester Vidal, SDM-TIB

Enrique Iglesias, University of Bonn

Diego Collarana, Fraunhofer IAIS

Oscar Corcho, OEG-UPM



DO NOT QUOTE ME ON THIS PRESENTATION

Thoughts of a junior researcher after ~4 years
working on Semantic Data Integration



“... not a precise, full methodology that others can use ...” (ISWC2019)

“... transformation process **should be** more general and less implementation-...” (ISWC2019)

“... focused on the use of standard formats ...” (ESWC2019)

“... I'm not convinced that it can be applied systematically ...”
(ESWC2019)

“... Although there was not much **scientific value**, I believe this paper to be of good quality and quite valuable as systems paper ...” (ESWC2020)

- 1) We were not doing **BASIC research**
- 2) Is this a **research problem?**

“I prefer to develop my own **ad-hoc script** for creating the KGC”

→ Ask Ana Iglesias about current status of the Bio2RDF

“Mappings are really **difficult to understand**”

create the mappings is very high”

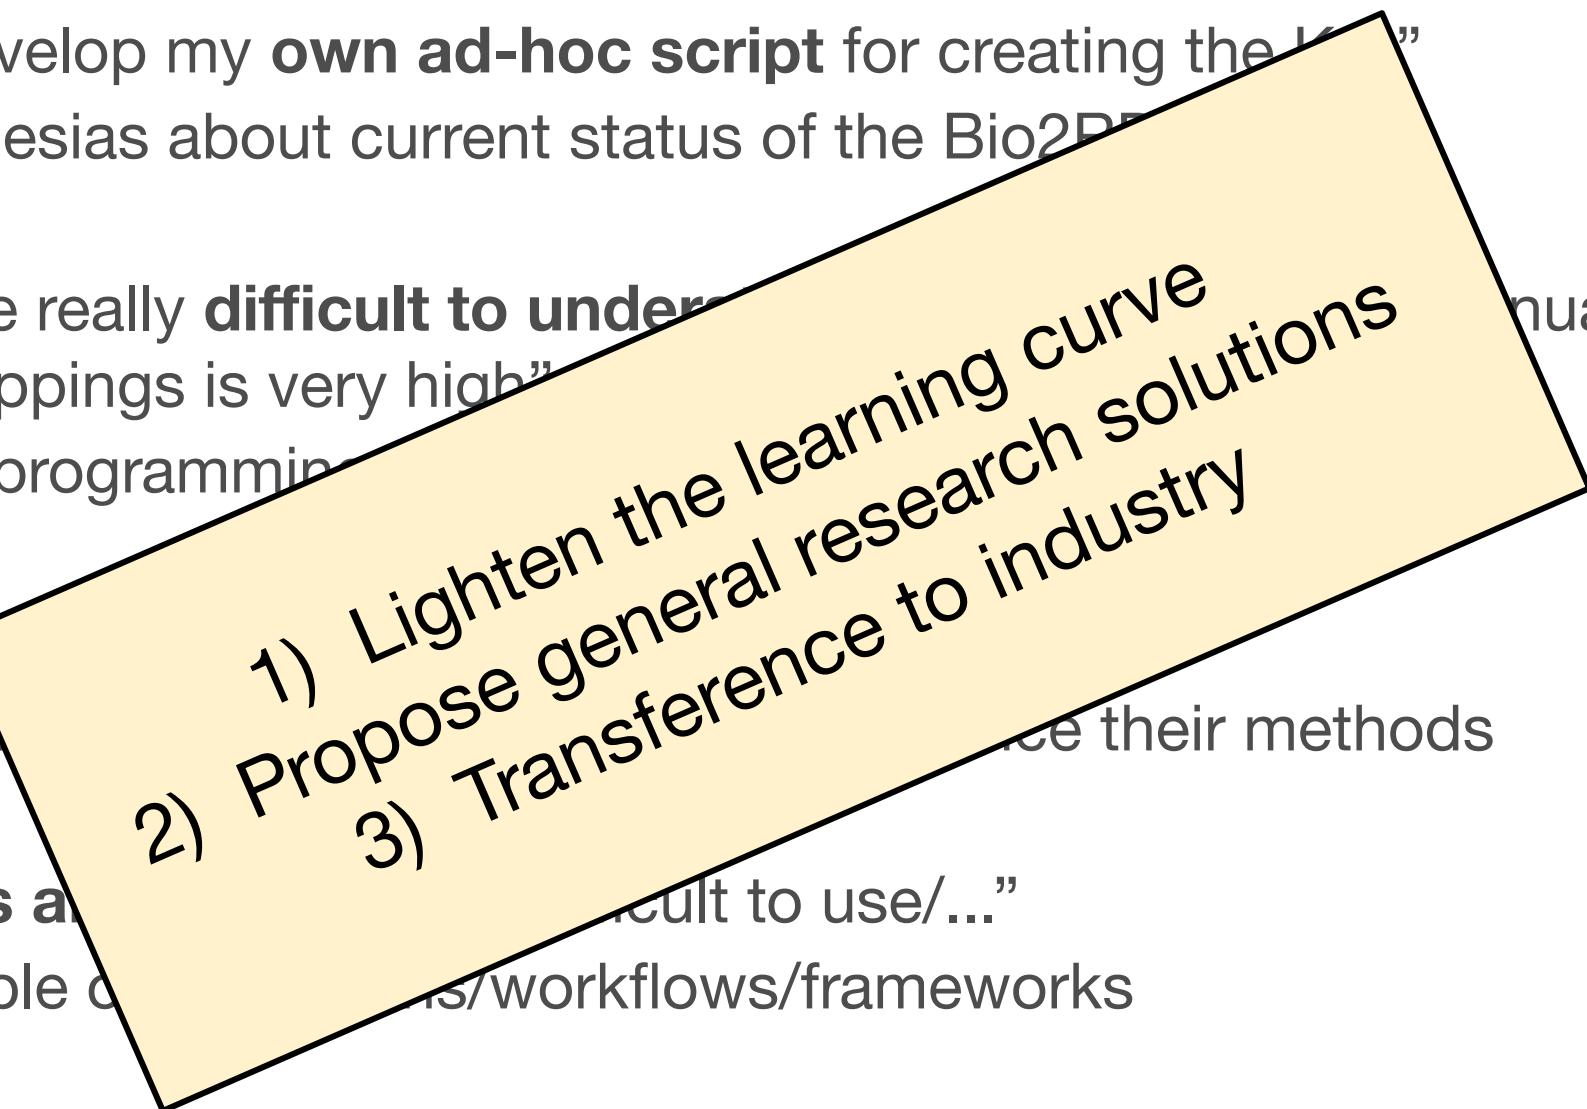
→ More than programming

“Your solution is...”

→ What is your solution?

“Your **engines** are...”

→ Generalizable or specific
to certain engines/workflows/frameworks

- 
- 1) Lighten the learning curve
 - 2) Propose general research solutions
 - 3) Transference to industry



- Inverting too much explaining your code
- Focus on the technology problems
- Specific solutions
- Try to explain ALL what I solved
- “Fight” with your competitors

The How

Not so important
for a research work

- Think before do
- Invest time in defining Research Questions
- Make your proposal general
- Implement the solution in a specific case(s)
- Discuss with your co-authors before doing anything
- Technology is a support not the core
- We are researchers, not engineers
- Good theoretical background of XXX (Databases)

The What

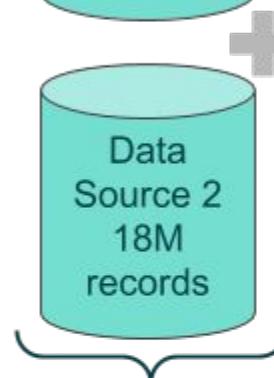
Not easy to find
and define

Enrique Iglesias, Samaneh Jozashoori*, David Chaves-Fraga*, Diego Collarana and Maria-Esther Vidal.* **SDM-RDFizer: An RML Interpreter for the Efficient Creation of RDF Knowledge Graphs.** Under review at CIKM20 Resource Track

Samaneh Jozashoori, David Chaves-Fraga*, Enrique Iglesias, Oscar Corcho and Maria-Esther Vidal.* **FunMap: Efficient Execution of Functional Mappings for Scaled-Up Knowledge Graph Creation.** Under review at ISWC20 Research Track

*The authors contributed equally to this research.

SDM-RDFizer: An RML Interpreter for the Efficient Creation of RDF Knowledge Graphs.



Total size ~ 1Gb
25% duplicates

```

1 <TriplesMap1>
2 rml:logicalSource [ rml:source "dataSource1" ];
3 rr:subjectMap [
4   rr:template "http://ialis.eu/{acc}_{enst}";
5   rr:class iasis:RBP_RNA_PhysicalInteraction ];
6
7 rr:predicateObjectMap [
8   rr:predicate iasis:interactionScore;
9   rr:objectMap [ rml:reference "omixcore" ] ];
10
11 rr:predicateObjectMap [
12   rr:predicate iasis:interaction_involves_RBP;
13   rr:objectMap [
14     rr:parentTriplesMap <TriplesMap2> ]].
15
16 <TriplesMap2>
17 rml:logicalSource [ rml:source "dataSource1" ];
18 rr:subjectMap [
19   rr:template "http://ialis.eu/Protein/{acc}";
20   rr:class iais:Protein ];
21
22 rr:predicateObjectMap [
23   rr:predicate iasis:protein_isRelatedTo_exon;
24   rr:objectMap [
25     rr:parentTriplesMap <TriplesMap3>;
26     rr:joinCondition [ rr:child "enst"; rr:parent "enst" ; ]];
27
28 <TriplesMap3>
29 rml:logicalSource [ rml:source "dataSource2" ];
30 rr:subjectMap [
31   rr:template "http://ialis.eu/Exon/{ense}";
32   rr:class iais:Exon ] .

```



RDF Knowledge Graph Creation



RocketRML :
Memory Failure



RMLMapper : TimeOut (>48h)



KG Embeddings

Question
Answering
Systems

Search Interfaces

Data Labelling

Explainable AI

Application Layer



Problem: Efficient knowledge graph creation in complex data integration scenarios

Objectives:

- O1) Define data structures that speed up the execution of mapping rules for KGC
- O2) Implement a set of unique physical operators for managing the data structures

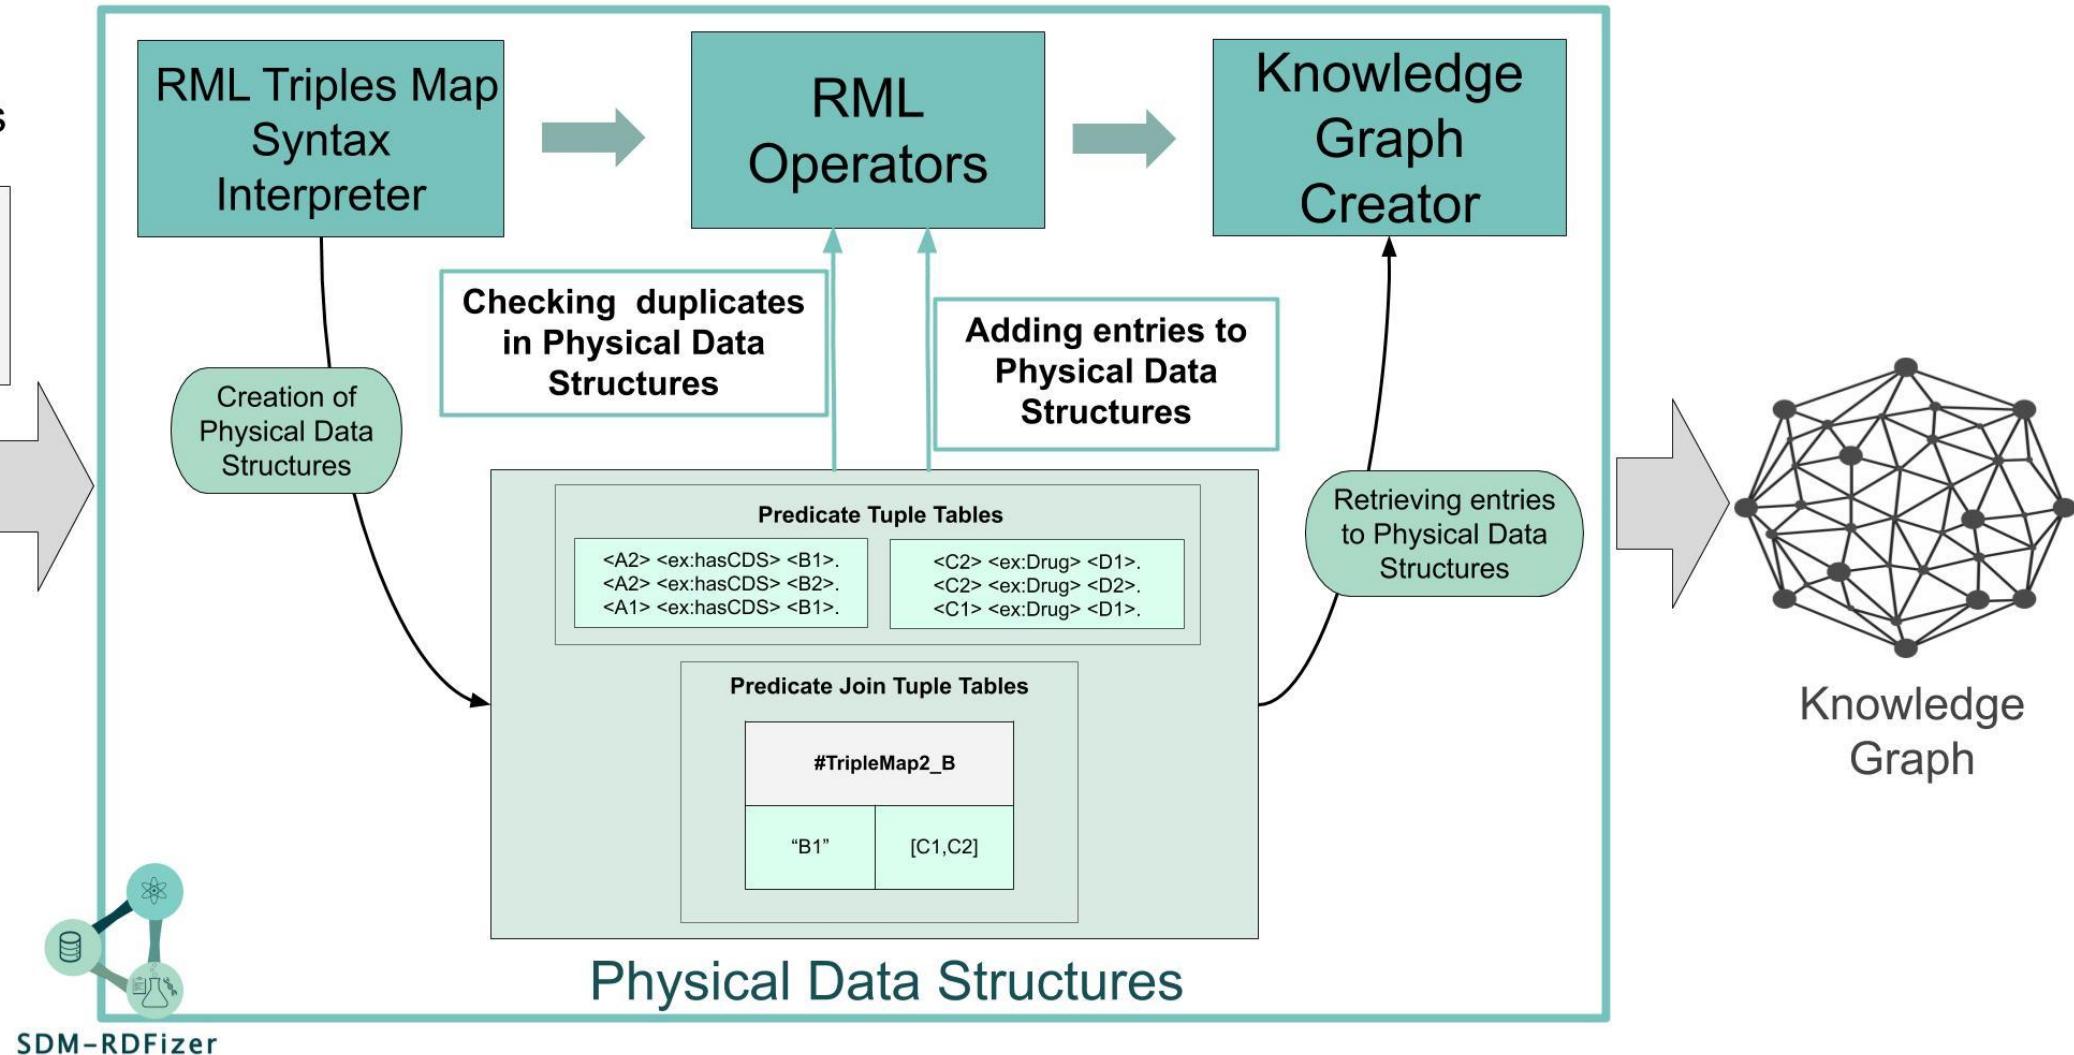
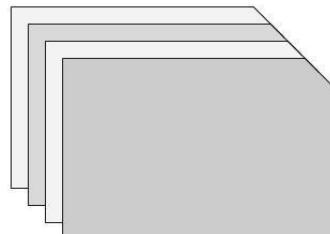
Research Questions:

- Q1) What is the impact of data duplication rate in the execution time of a knowledge graph creation approach?
- Q2) What is the impact of input data size in the total execution time of a knowledge graph creation process?
- Q3) What is the effect of the triples map types in the PredicateObjectMap over the existing engines?

RML Triples Maps

```
<#TripleMap1>
rr:subjectMap [
  rr:template "{A}";
  rr:objectMap [
    rr:predicatex:hasCDS;
    rr:objectMap [
      rml:reference "{B}";
    ];
  ];
];
```

Data Sources



SDM-RDFizer



Number of operations for KGC	Naïve Approach	SDM-RDFizer operators
Simple Object Map	$ N_p + S_p + \Theta(N_p \log(N_p))$	$ N_p + 2 S_p $
Object Reference Map (natural join)	$ N_p + S_p + \Theta(N_p \log(N_p))$	$ N_p + 2 S_p $
Object Join Map (join)	$ N_{parent} \times N_{child} + N_p + S_p + \Theta(N_p \log(N_p))$	$2 N_{parent} + N_{child} + N_p + 2 S_p $

$|N_p|$: Cardinality of a multiset for all RDF triples of p, where p is a predicate (RDF KG with duplicates)

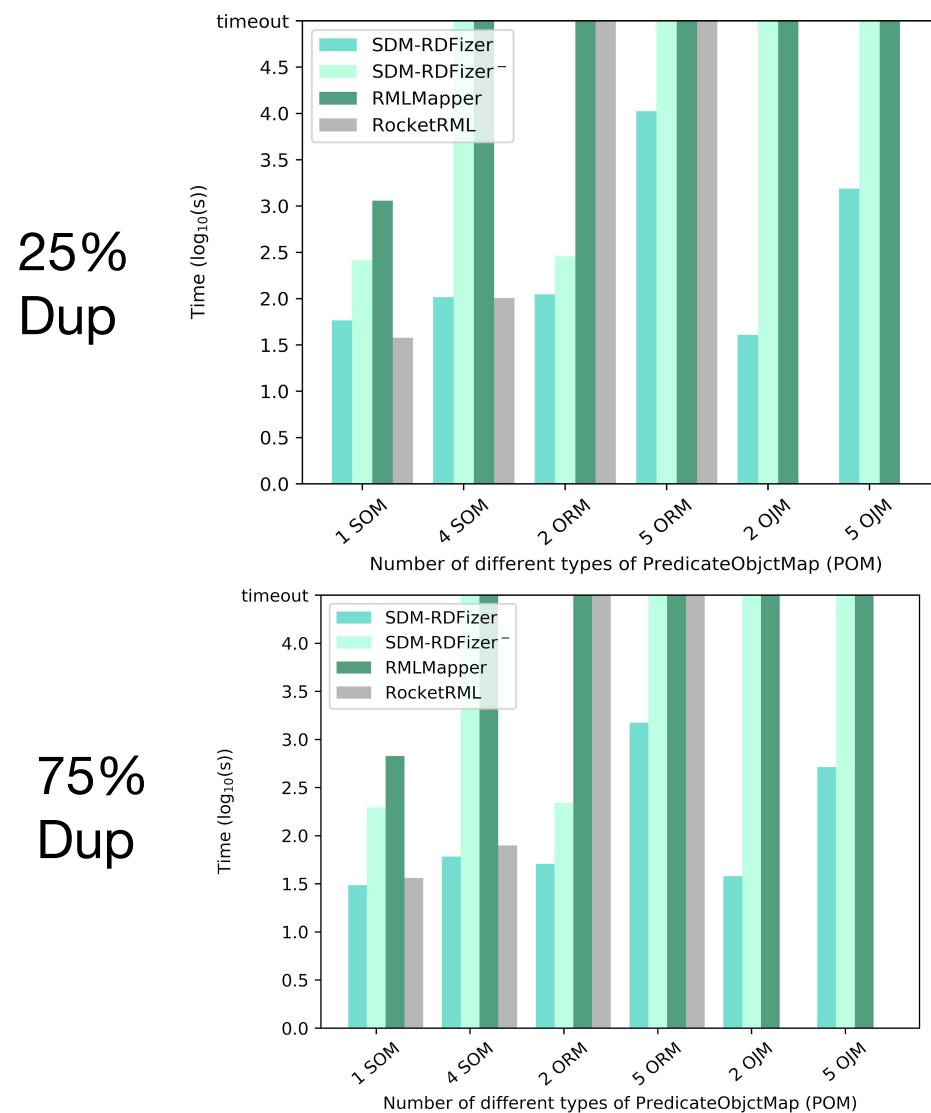
$|S_p|$: Cardinality of a set all RDF triples of p, where p is a predicate (RDF KG without duplicates)

$\Theta(\cdot)$: Duplicates removal algorithm $|N_{parent/child}|$: Columns in parent/child

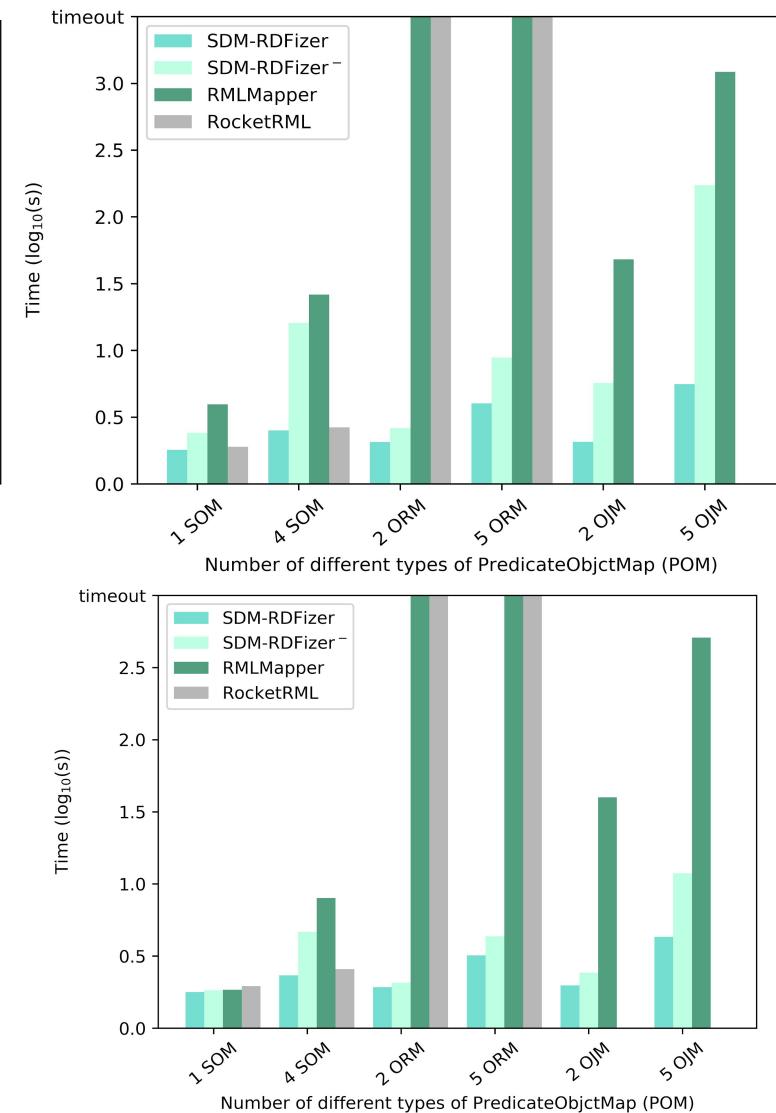
Experimental results



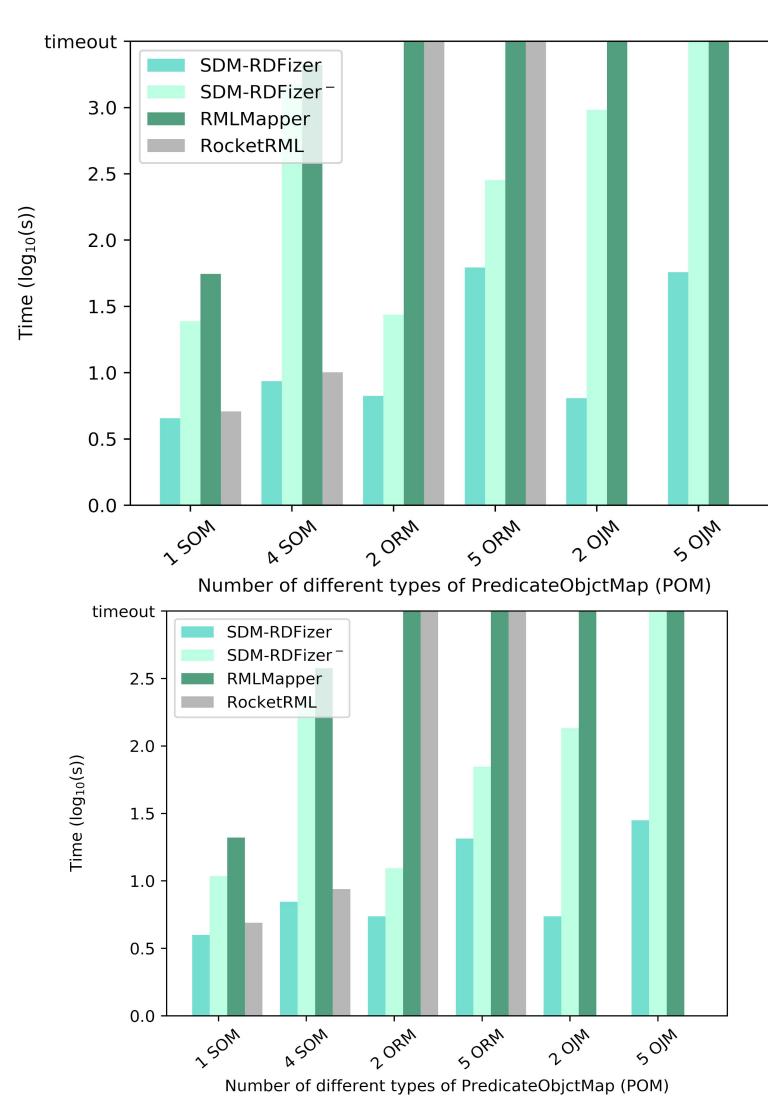
10K rows



100K rows

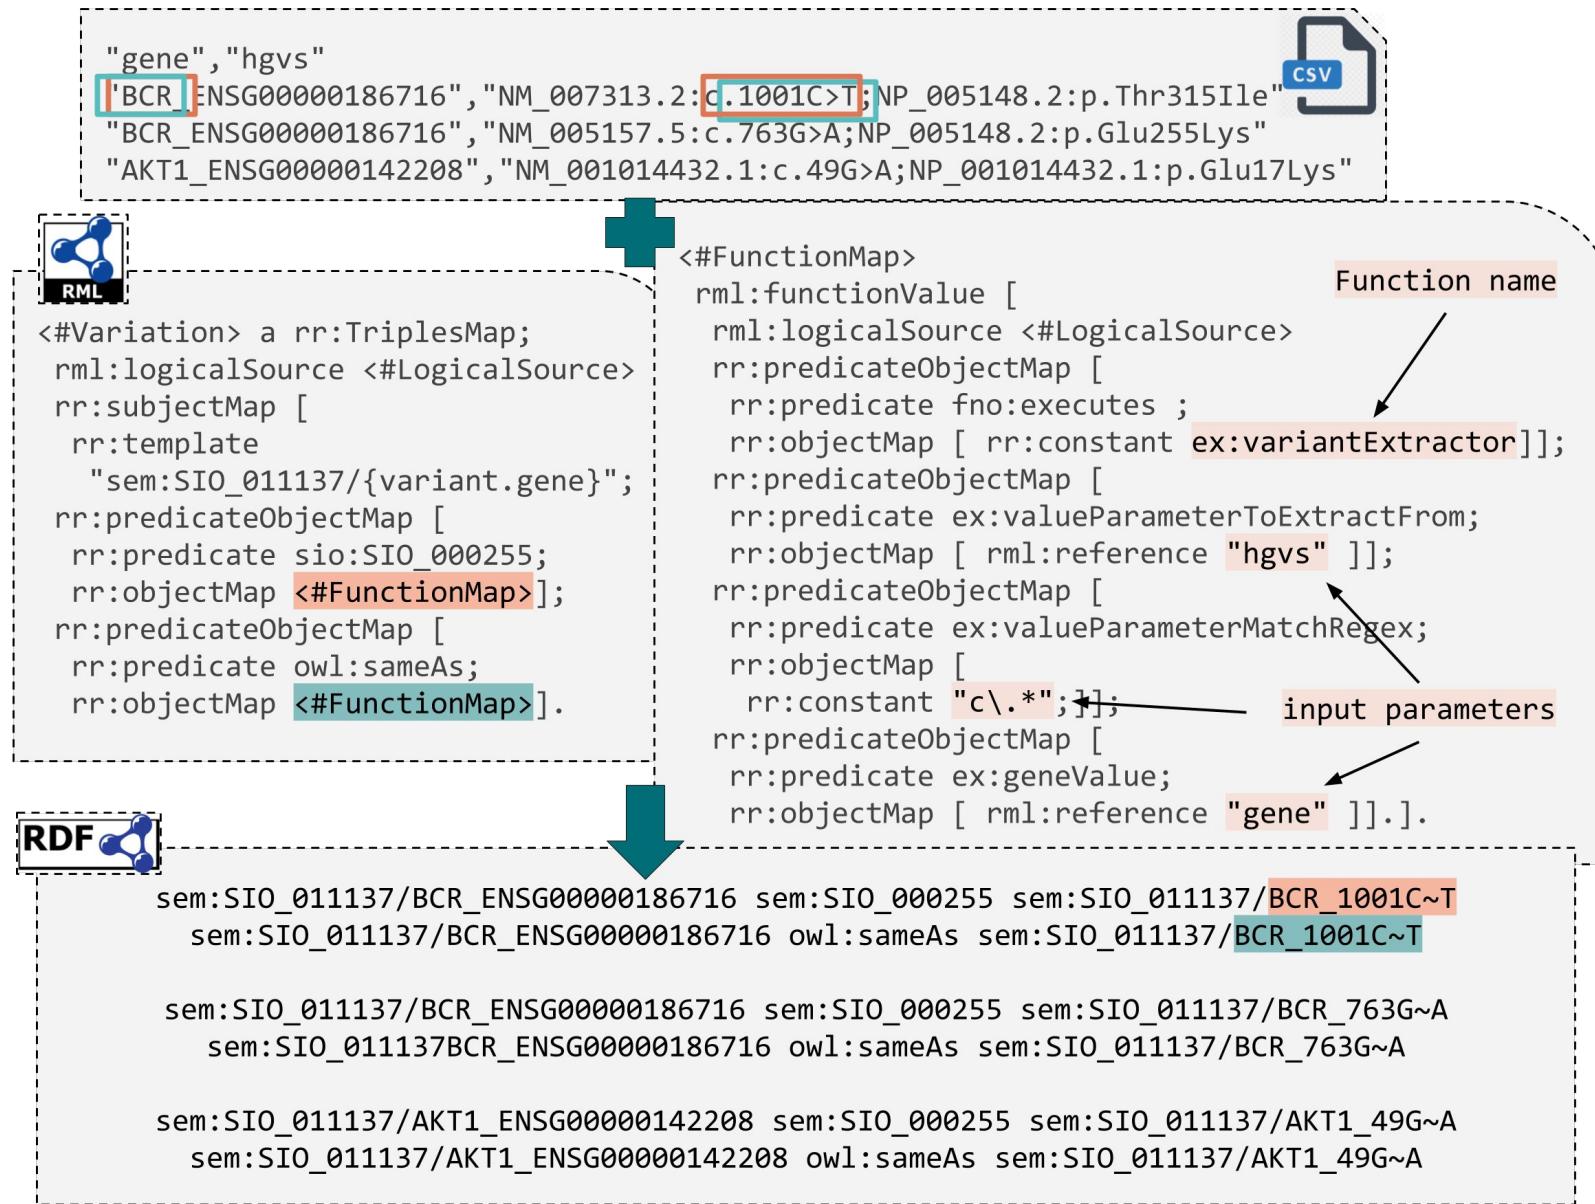


1M rows



- **Novel physical operators and data structures** that speed up the generation of duplicate-free KG
- Empirical results indicate that SDM-RDFizer **outperforms the state of the art** by up to three orders of magnitude
- Basis for the development of **real-world knowledge graph applications**

FunMap: Efficient Execution of Functional Mappings for Scaled-Up Knowledge Graph Creation





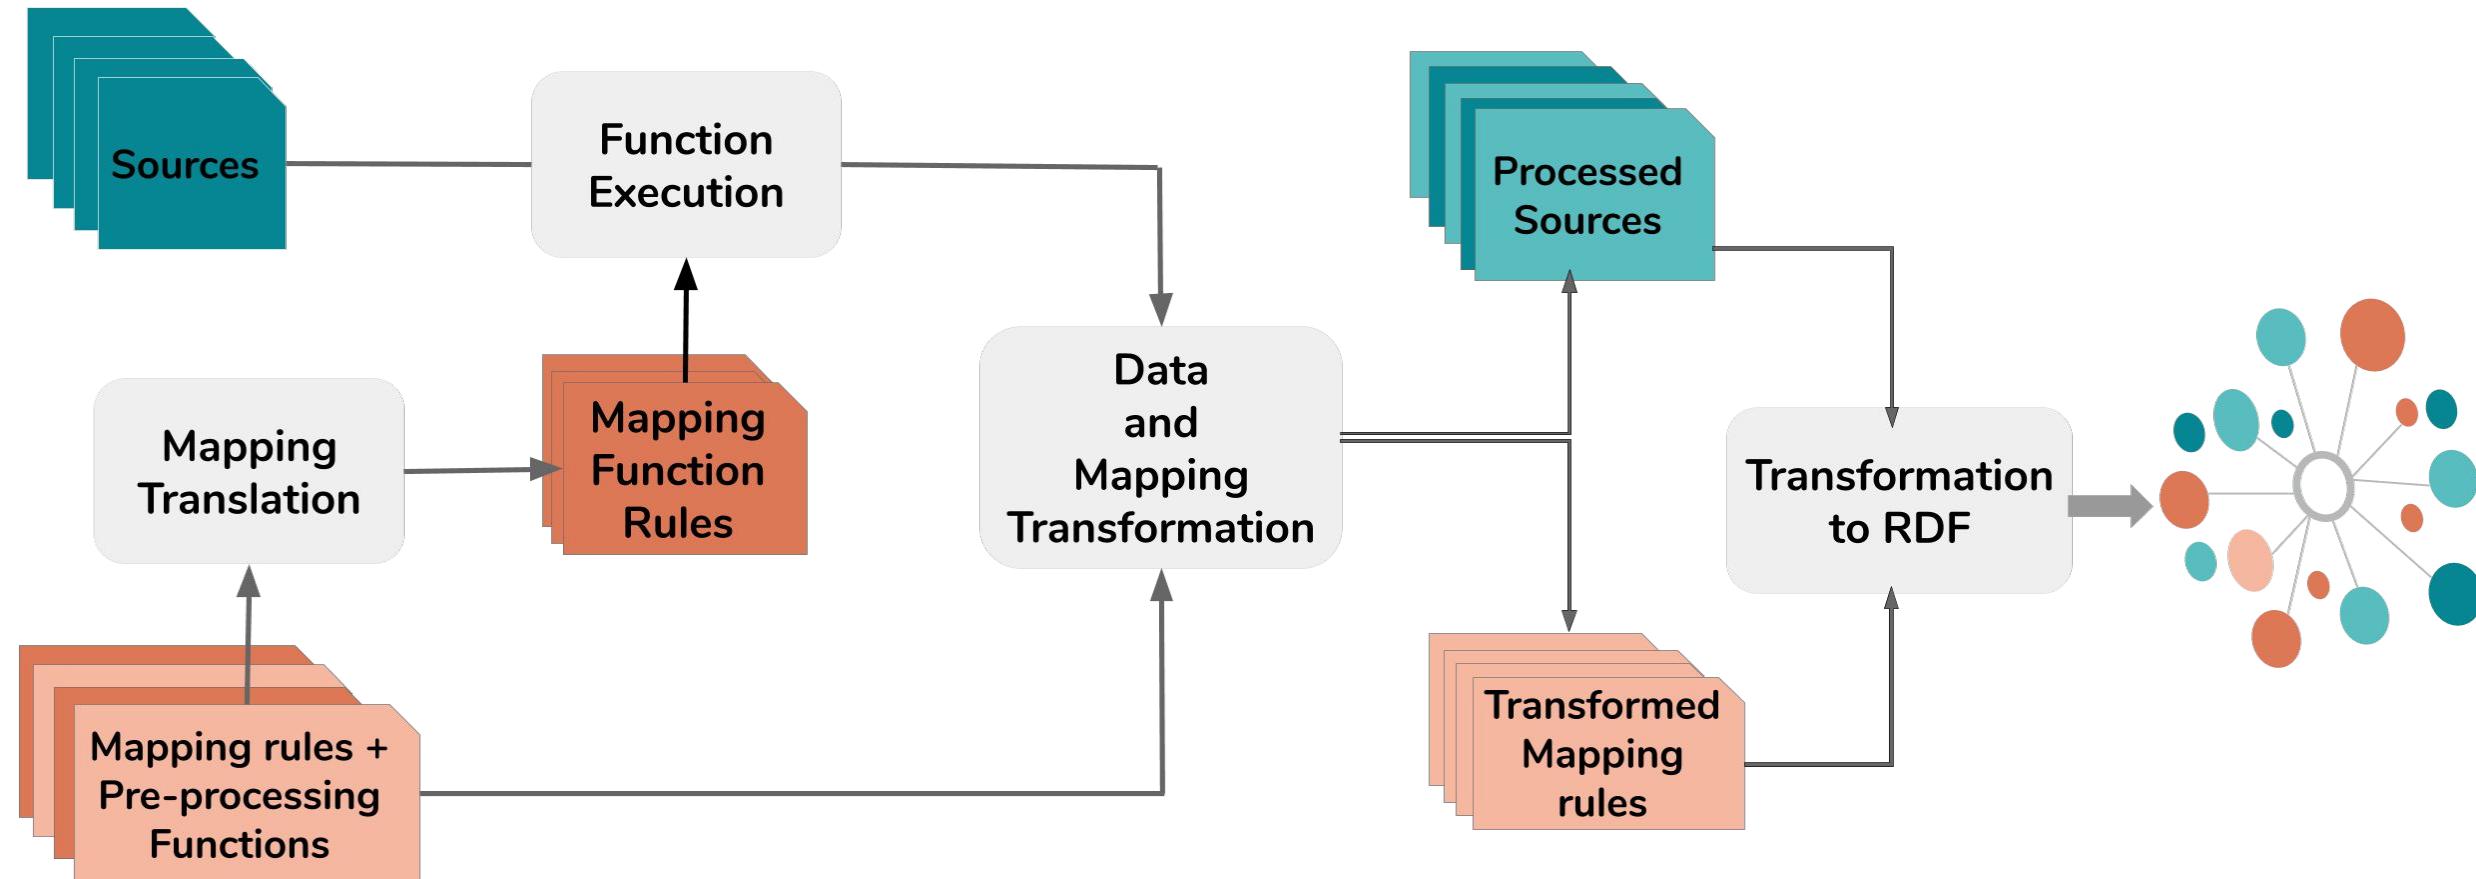
Problem: Scaled-up KG construction from functional mapping rules

Objectives:

- O1) Transform a data integration system with functional mappings into an equivalent data integration system where mappings are function-free
- O2) Optimization techniques to reduce the total execution time of the KGC

Research Questions:

- Q1) What is the impact of data duplication rate in the execution time of a knowledge graph creation approach?
- Q2) What is the impact of different types of complexity over transformation functions during a knowledge graph creation process?
- Q3) How does the repetition of a same function in different mappings affect the existing RML engines?







ID	Mutation genome position	GENOMIC_MUTATION_ID	ID	Mutation genome position	Gene name	ID	Mutation genome position	functionOutput
1	22:20302597-20302597	COSV50619134	1	22:20302597-20302597	DGCR6L	1	22:20302597-20302597	22:20302597:20302597
3	17:18874996-18874996	COSV58755801	2	1:186072702-186072702	HMCN1	3	17:18874996-18874996	17:18874996:18874996
4	1:186072702-186072702	COSV54901969	3	17:18874996-18874996	SLC5A10_ET0000039564	4	1:186072702-186072702	1:186072702:186072702
5	6:56246049-56246049	COSV63690608	4	1:186072702-186072702	HMCN1_ET00000367492	5	6:56246049-56246049	6:56246049:56246049
6	1:243692781-243692781	COSV55606438	5	6:56246049-56246049	COL21A1_ET0000037081	6	1:243692781-243692781	1:243692781:243692781
7	10:50044166-50044166	COSV55433638	6	1:243692781-243692781	AKT3	7	10:50044166-50044166	10:50044166:50044166
...	7	10:50044166-50044166	WDFY4_ET00000413659

Projected 1

Projected 2

Output 1

```

<#TriplesMap1>
  rml:logicalSource [ rml:source "source1.csv";
                      rml:referenceFormulation ql:CSV ];
  rr:subjectMap <#FunctionMap1>;
  rr:predicateObjectMap [
    rr:predicate iasis:represents;
    rr:objectMap [
      rml:reference "Mutation" ];
  rr:predicateObjectMap [
    rr:predicate iasis:tissue;
    rr:objectMap [
      rml:reference "Primary site" ]].

```

```

<#FunctionMap1>
  a fnml:FunctionTermMap;
  fnml:functionValue [
    rml:logicalSource [ rml:source "source1.csv";
                        rml:referenceFormulation ql:CSV ];
    rr:predicateObjectMap [
      rr:predicate fno:executes ;
      rr:objectMap [
        rr:constant ex:replaceValue ];
    rr:predicateObjectMap [
      rr:predicate ex:value;
      rr:objectMap [
        rml:reference "Mutation genome position"]];
    rr:predicateObjectMap [
      rr:predicate ex:value2;
      rr:objectMap [
        rr:constant "-" ];
    rr:predicateObjectMap [
      rr:predicate ex:value3;
      rr:objectMap [
        rr:constant ":" ]];
  
```

Transforms to

```

<#TriplesMap1>
  rml:logicalSource [ rml:source "output1.csv";
                      rml:referenceFormulation ql:CSV ];
  rr:subjectMap [
    rml:reference "functionOutput" ];
  rr:predicateObjectMap [
    rr:predicate iasis:represents;
    rr:objectMap [
      rr:parentTriplesMap <#TriplesMap2>;
    rr:joinCondition [
      rr:child "Mutation genome position";
      rr:parent "Mutation genome position"];
    ];
  rr:predicateObjectMap [
    rr:predicate iasis:tissue;
    rr:objectMap [
      rr:parentTriplesMap <#TriplesMap3>;
    rr:joinCondition [
      rr:child "Mutation genome position";
      rr:parent "Mutation genome position"];
    ];
  ];

```

```

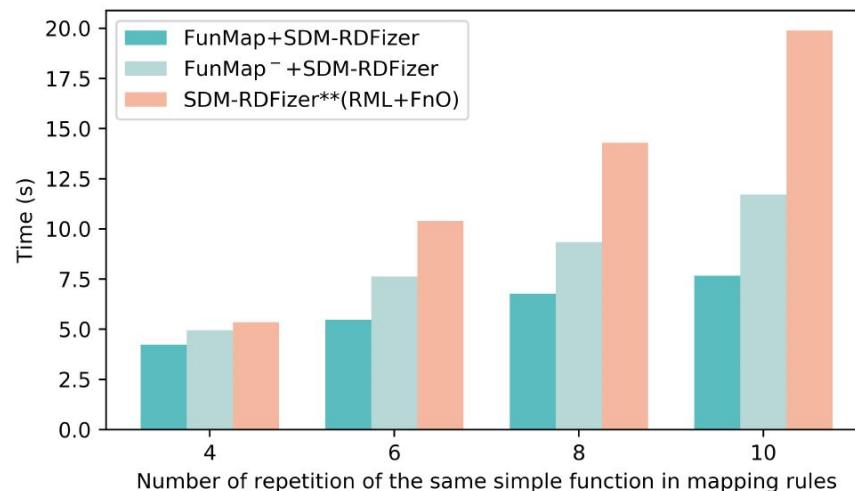
<#TriplesMap2>
  a rr:TriplesMap;
  rml:logicalSource [ rml:source "projected1.csv";
                      rml:referenceFormulation ql:CSV ];
  rr:subjectMap [
    rml:reference "Mutation" ].

<#TriplesMap3>
  a rr:TriplesMap;
  rml:logicalSource [ rml:source "projected1.csv";
                      rml:referenceFormulation ql:CSV ];
  rr:subjectMap [
    rml:reference "Primary site" ].

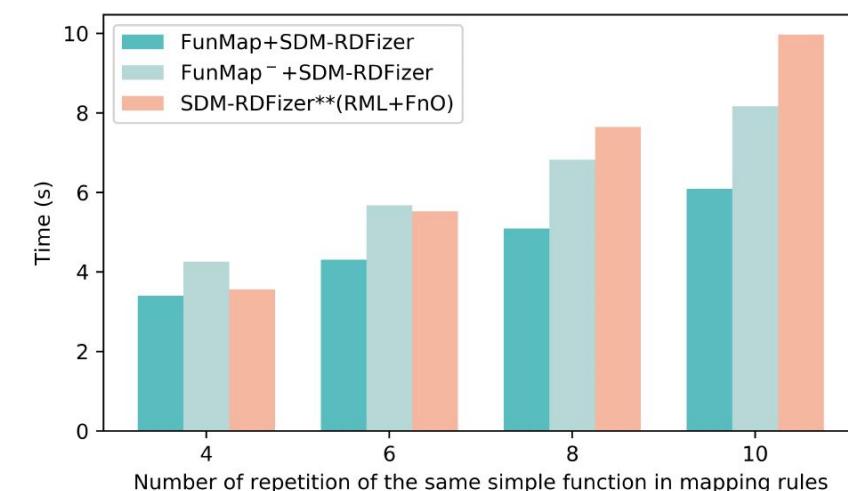
```

FunMap over SDM-RDFizer - Experimental results

Simple functions
(lower, upper)

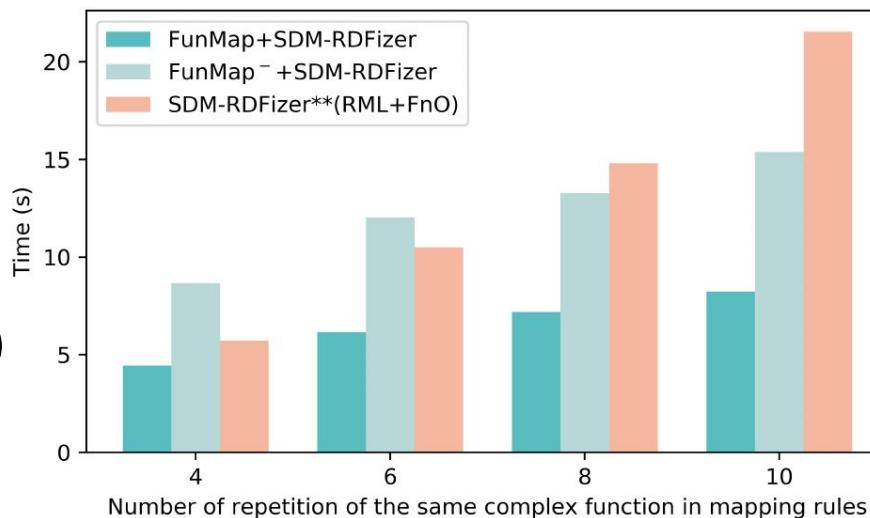


(a) SDM-RDFizer - 25% of duplicates

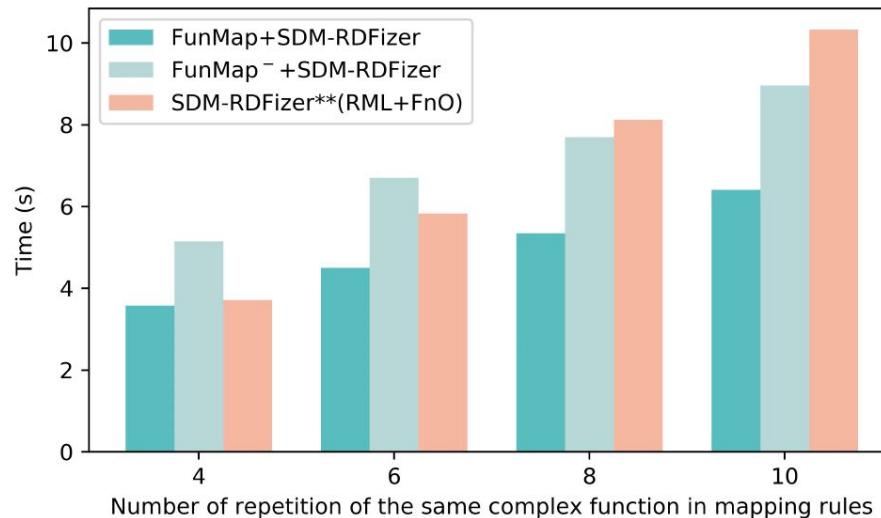


(b) SDM-RDFizer - 75% of duplicates

Complex functions
(if, replace,
multiple columns)



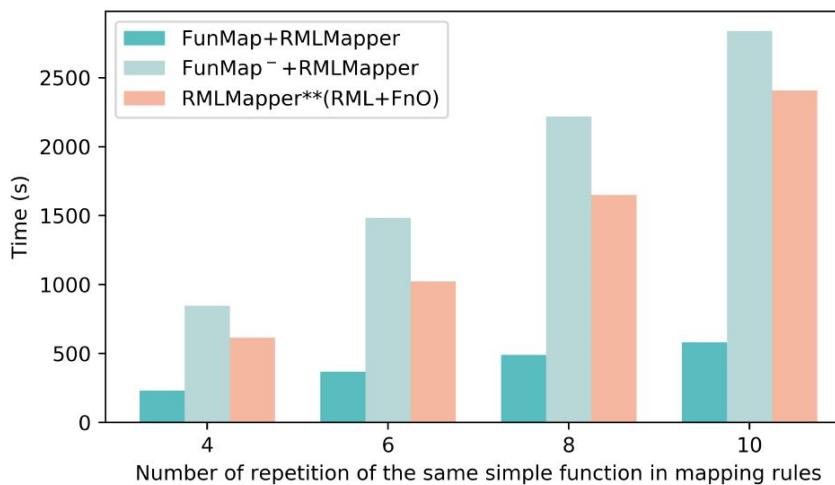
(a) SDM-RDFizer - 25% of duplicates



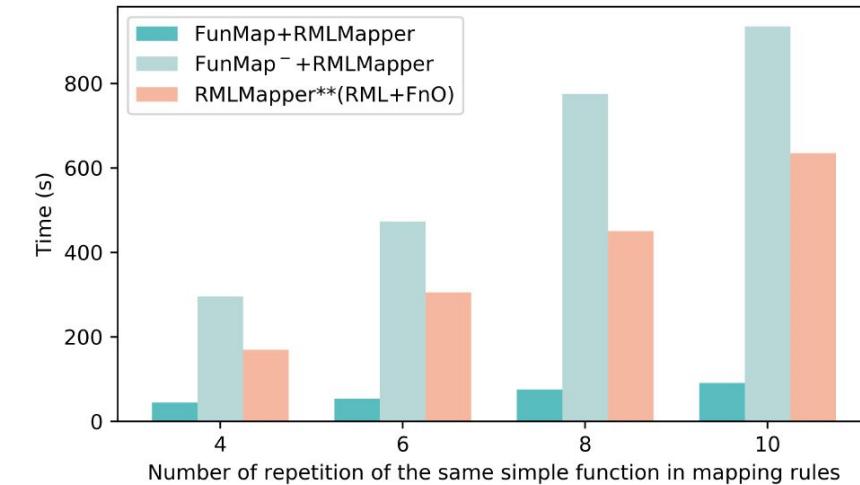
(b) SDM-RDFizer - 75% of duplicates

FunMap over RMLMapper - Experimental results

Simple functions
(lower, upper)

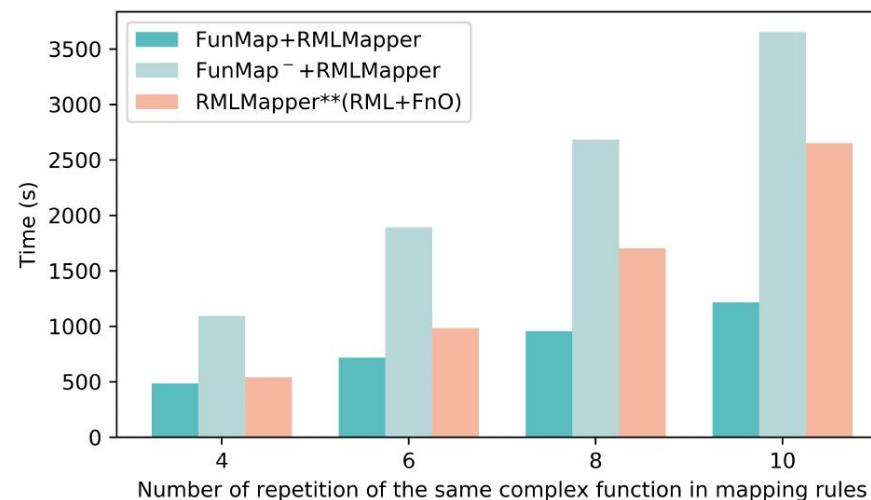


(c) RMLMapper - 25% of duplicates

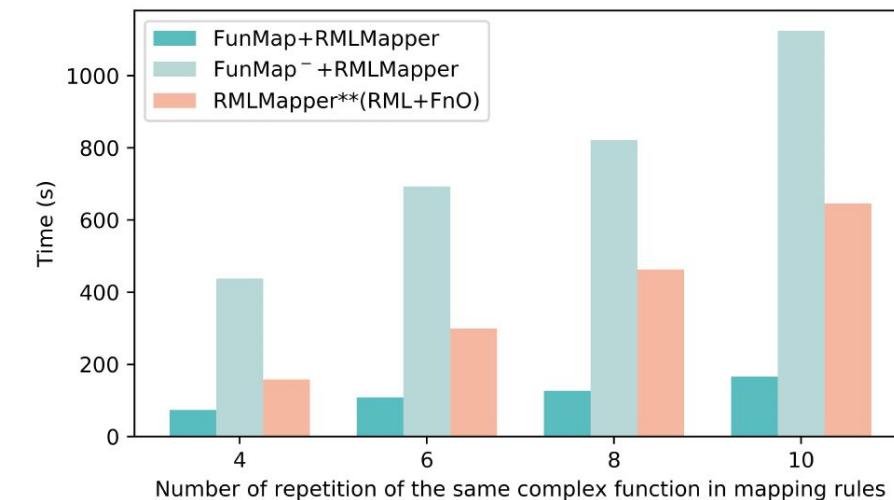


(d) RMLMapper - 75% of duplicates

Complex functions
(if, replace,
multiple columns)



(c) RMLMapper - 25% of duplicates

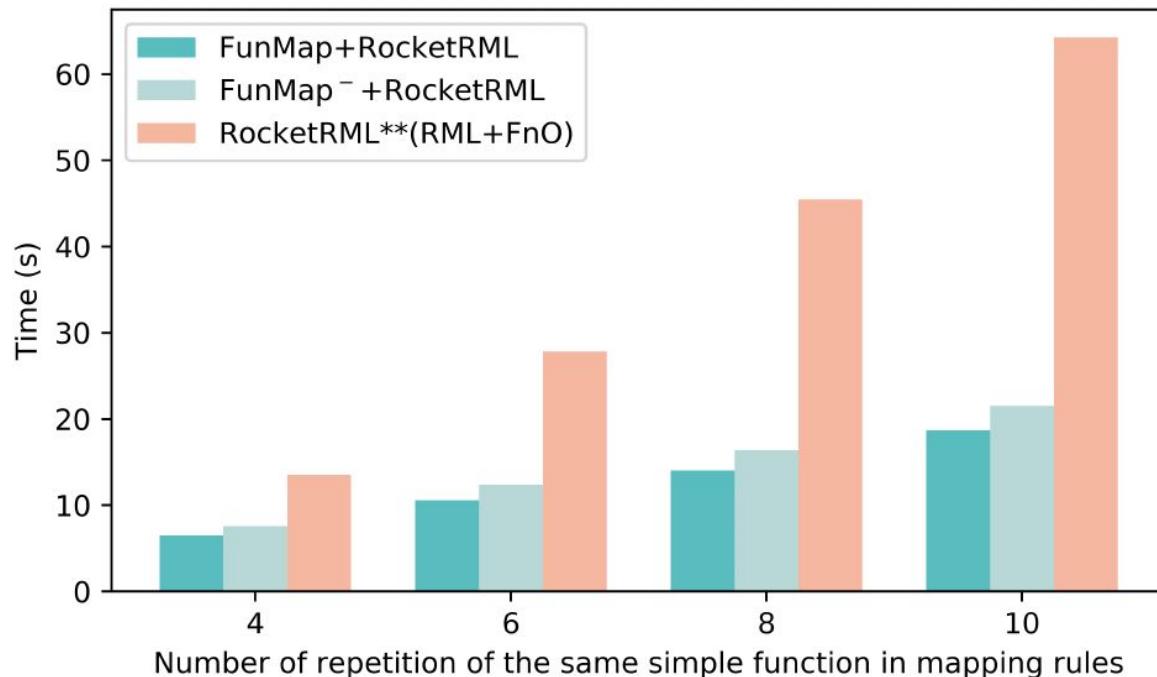


(d) RMLMapper - 75% of duplicates

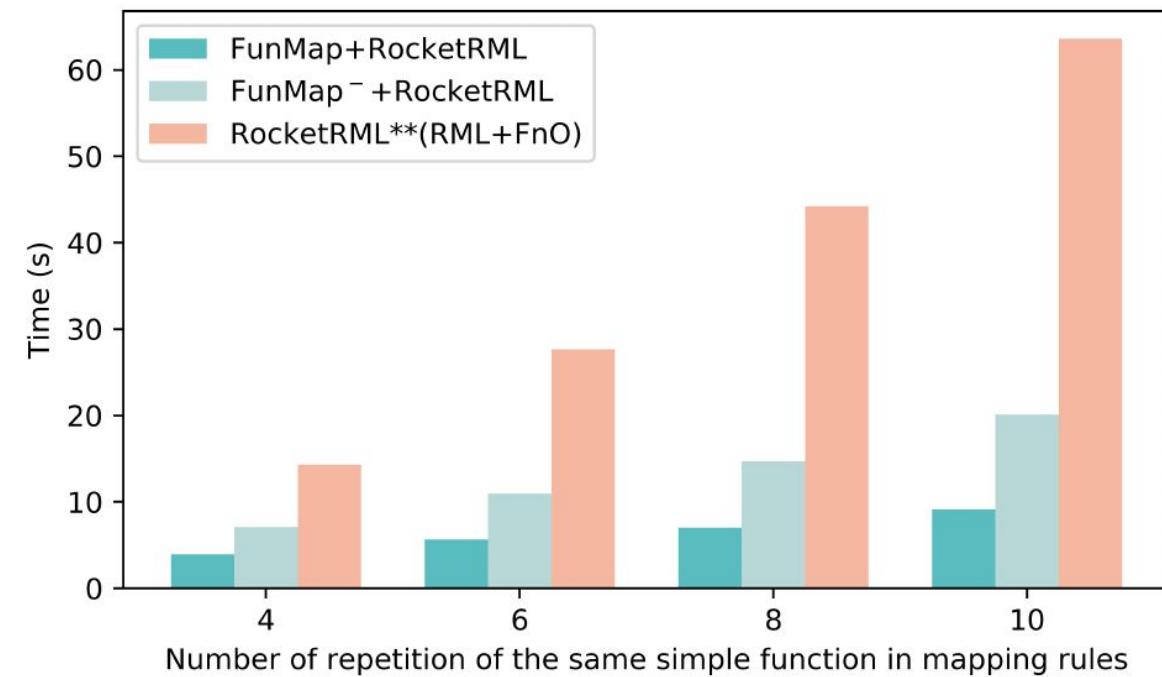


FunMap over RocketRML - Experimental results

Simple functions (lower, upper)



(e) RocketRML - 25% of duplicates



(f) RocketRML - 75% of duplicates

- Heuristic-based approach for **generating scalable data integration systems**
- FunMap converts data integration systems in RML+FnO into **equivalent data integration systems specified in RML**
- FunMap generates data integration systems that **enhance RML-complaint engines**
- Empirical evaluations suggest that the execution time of RML+FnO **is reduced by up to 20 times**



- Research takes time:
 - SDM-RDFizer: March 2019 - June 2020
 - FunMap: October 2019 - May 2020
- (Try to) be on the same page with your co-authors/supervisors
- Be passionate and believe in what you do
- Envision big/general/global and start small/specific/local
- Be patient with (Semantic Web) reviewers
- Your impact will be as big as the quality of your pitch/paper to explain the solution*

*Pieter Colpaert at Open Summer of Code



From Engineering To Research

The KG Construction Use Case

David Chaves-Fraga, Ontology Engineering Group
Universidad Politécnica de Madrid, Spain

Samaneh Jozashoori, SDM-TIB

Maria-Ester Vidal, SDM-TIB

Enrique Iglesias, University of Bonn

Diego Collarana, Fraunhofer IAIS

Oscar Corcho, OEG-UPM

